

ARM26BX06-NV012 / ARM26TX06-NV003 - Agentic-AI, Schema-Driven Decision Management

Army (DEVCOM GVSC) | SBIR (NV012) and STTR twin (NV003) | Release 6 CSO | Opens Sep 23, 2026 (TPOC direct contact ends) | DSIP Q&A closes Oct 7, 12:00 ET | Proposals due Oct 21, 12:00 ET

Recommendation: PURSUE (submit to one twin, not both). The Q&A removed the two things that made the topic look wired - the DAOSoft dependency and the need for Army data. Phase I is 'proof by demonstration', synthetic data is explicitly acceptable, CMMC Level 1, no ITAR. Confidence HIGH on fit.

Phase I award	Up to \$300,000, 1-6 months (Army CSO, 'Phase I (Feasibility)')
Proposal format	Army CSO technical volume: 7 pages for Phase I (verify against Section 3 of the Army instructions)
Compliance	CMMC Level 1. No ITAR clause. Army CSO (Commercial Solutions Opening) evaluation, not classic BAA.
Phase II	Two cases from PAE Maneuver Ground and PAE Fires; COTS-style product the Army and industry would both use
STTR twin	NV003: identical text; needs a research institution at >= 30% of the work

What they want

[TOPIC TEXT] "Prove a decision-as-a-program capability that turns complex engineering/acquisition questions into structured, auditable Decision Packages - faster, with higher confidence."

Phase I desired outcomes:

- A decision schema in which objectives, options, constraints, assumptions, risks and bias checks are first-class objects.
- An agentic layer that performs structured elicitation into that schema, generates a decision workflow plan, and runs reproducible evaluations with explicit "what flips the decision" logic.
- Two end-to-end demonstrations: one point-in-time trade study, one long-horizon decision program with refresh cycles.

Context engineering for the ground-vehicle domain (knowledge graphs / semantic layers), interoperability with digital-engineering artifacts (SysML 2.0, UML), and human control per human-AI interaction guidance are all named.

What the Q&A changed

- [Q&A] "DAOSoft is not a requirement!" - offerors may use their own platform, integrate with DAOSoft, or replace it.
- [Q&A] "For Phase I, offeror generic documents are sufficient" and "offeror synthetic modeling is sufficient if PAE or PM data is not forthcoming" - no government data wall.
- [Q&A] Either demonstration structure is acceptable; "the government supports a fresh innovative approach."
- [Q&A] On 'open-source' in Phase II: "COTS is better... Offerors will own their own code, and the government places no restrictions on an offeror chosen distribution model." The intent is avoiding government-maintained code.
- [Q&A] The effort is framed as "100% new effort" toward agentic AI for rapid product decisions.

How we would build it

- **Schema:** a typed decision model (JSON Schema / Pydantic) - Decision, Option, Criterion, Constraint, Assumption (with evidence and confidence), Risk, BiasCheck, RefreshTrigger - versioned so a decision program is a sequence of schema instances over time.
- **Elicitation agent:** LLM agent that interviews a user and ingests documents (requirements, spec sheets, cost tables) into the schema with citations; every field carries provenance.
- **Evaluation engine:** deterministic scoring (weighted criteria, constraint satisfaction, Monte Carlo on uncertain assumptions) producing the ranked options and a sensitivity table: which assumption or weight flips the top choice.
- **Refresh loop:** when an assumption's evidence changes, re-run and produce a delta report - this is the 'long-horizon decision program' demonstration.
- **Audit/trust layer:** bias checks as executable tests (anchoring on incumbent option, missing alternatives, unstated assumptions), full run logs, signer-ready package export (PDF/SysML-compatible artifact).
- **Demonstrations:** (1) a vehicle-subsystem trade study built from public sources (e.g., powertrain or armor options with published cost/weight data); (2) a multi-year 'program' where new evidence arrives each cycle.

Resources we would use

- Public ground-vehicle trade-study material (SAE papers, GAO/CRS reports on Army vehicle programs) for synthetic cases.
- SysML v2 open tooling (pilot implementation / API) to show interoperability with digital-engineering artifacts.
- The topic's own reference: SAE 2025-01-0455 (agentic decision-intelligence framework) - align vocabulary with the author.
- DAOSoft public site (daosoft.ai) to understand the incumbent's structured-decision UX without depending on it.
- LLM agent framework of choice; graph store for the semantic layer if needed. No hardware.

Bottlenecks and risks

- **Incumbent adjacency:** DAOSoft is on GVSC's network and the topic author knows it; a proposal must show clear advance beyond 'structured trade study + LLM', i.e., reproducibility, sensitivity and refresh logic.
- **Acquisition-domain credibility:** Army acquisition decision programs have their own vocabulary (AoA, KPPs, milestone decisions). A short engagement with a former acquisition professional would de-risk the demos.
- **Determinism vs. LLMs:** reviewers will ask how an LLM-driven pipeline is 'reproducible'; the architecture must confine the LLM to elicitation/authoring and keep evaluation deterministic.
- **Page limit:** 7 pages is very short for a two-demo plan; the schema and the sensitivity logic must be shown compactly (one figure each).
- **SBIR vs STTR choice:** the STTR twin requires an RI partner; only worth it if a university decision-science or MBSE group adds credibility we cannot show ourselves.

Decision factors

- Fit to team: schema + workflow + evaluation harness + LLM orchestration - direct.
- Dual use: [TOPIC TEXT] "the vehicle design process for the automotive industry"; more broadly any product-development trade-study workflow (a generic 'decision package' tool).
- Phase III framing is ERP-style deployment per customer - a services-heavy model; decide whether that is a business you want.
- Decisive pre-proposal experiment: encode one public trade-off in the schema, run the deterministic evaluator, then perturb one assumption and confirm the 'what flips the decision' output is correct and traceable.

Before Oct 7 (Q&A): the PAE data answers are still 'pending' - ask whether any unclassified PAE Maneuver Ground case will be released in Phase I, and whether a SysML v2 export is expected or merely desirable.

Tags: [TOPIC TEXT] = quoted from the solicitation; [Q&A] = government answer on DSIP; [INFERRED] = our judgment, not verified by search. Generated 2026-09-20 from release6_triage.md / triage_results.md / sitis_qa.md.